

Scenario

DataTek, a customer analytics company, has developed a working churn prediction solution (the local codebase you were given). The leadership team now wants to **operationalize this model at scale** so business teams can consume predictions reliably through an API. They also require the solution to be production-grade with versioning, automated retraining, CI/CD, and monitoring, all aligned to **Azure Machine Learning best practices**.

Your task is to transform the existing repo into a full **Azure MLOps pipeline**.

Code is provided (data_utils.py, train.py, predict.py)

Core Requirements

1. Model Training Pipeline

- Adapt the training script so that it runs as an **Azure ML job** on a managed compute cluster.
- Ensure the job accepts input dataset parameters and outputs a registered model to the **Azure ML Model Registry**.

2. Model Deployment

- Deploy the registered model as an **Azure ML Managed Online Endpoint**.
- The endpoint should expose a REST API for real-time churn predictions.
- Ensure deployment follows best practices (staging vs. production deployments, traffic splitting, rollbacks).

3. CI/CD Integration

- Use **GitHub Actions** to build an automated pipeline that:
 - Runs unit tests on every PR.
 - Submits training jobs to Azure ML.
 - On success, registers the new model in the registry.
 - Deploys the model automatically to a **staging endpoint**.
 - Provides a manual approval gate before promoting the model to production.

4. Monitoring & Alerts

- Integrate **Application Insights** or **Azure Monitor** with the online endpoint.

- Track prediction latencies, error rates, and data drift (feature distributions vs. baseline).
- Provide an alerting plan (e.g., email or Teams notifications) if drift or high error rates are detected.

5. Governance & Security

- Store sensitive configs (keys, connection strings) in **Azure Key Vault**.
- Document a strategy for safe rollback in case of faulty deployments.

Deliverables

- **Technical Report/PPT (Max 5 slides)** describing:
 - Overall MLOps architecture on Azure ML.
 - Key design decisions (compute, environments, monitoring).
 - Cost optimization and governance/security considerations.
- **GitHub Repository** containing:
 - Modularized training and scoring scripts (parameterized).
 - GitHub Actions workflow (.github/workflows/mlops.yml).
 - Azure ML job definitions (YAML or Python SDK).
 - Unit tests for training, preprocessing, and inference functions.
 - Instructions in README.md for end-to-end execution.
- **Deployed Online Endpoint:**
 - API endpoint showing prediction response.
- **Monitoring Plan:**
 - Evidence of integrated Application Insights/monitoring/Drift Detection